

Compiling Hadoop after making a change in its code

1. After making a change to java class/classes in hadoop-common-project (or any module), execute `mvn compile -Pnative` from the module.

For example, in my computer, the project is at the location: `~/Documents/Our_Project/hadoop-3.0.0-src/`

```
hduser@lakshmijm:~/Documents/Our_Project/hadoop-3.0.0-src/hadoop-common-project$ mvn compile -Pnative
```

Wait till the compilation finishes successfully. If there are errors, fix them and recompile.

2. Then, go to parent project and execute

```
hduser@lakshmijm:~/Documents/Our_Project/hadoop-3.0.0-src$ sudo mvn package -Pnative -Pdist -Dtar -DskipTests -Drequire.isal -Disal.lib=/usr/lib -Dbundle.isal
```

This is for creating the tar of the new project's distribution which is to be deployed on to all instances on our Oracle cluster.

(Note: Flags for native code and ISA-L features are on in the mvn build command)

Wait for the completion of successful build of all project modules.

The newly built project's tar would be available at the following location in the file `~/Documents/Our_Project/hadoop-3.0.0-src/hadoop-dist/target/hadoop-3.0.0.tar.gz`

3. Once this is done, we can run our bash script `mvn_project_configure.sh` to configure our nodes with the latest version of the Hadoop project.

(Important step in the script: Copy the configuration files of our Oracle cluster into the location `$HADOOP_HOME/etc/hadoop`)